


ARTICLE OPEN

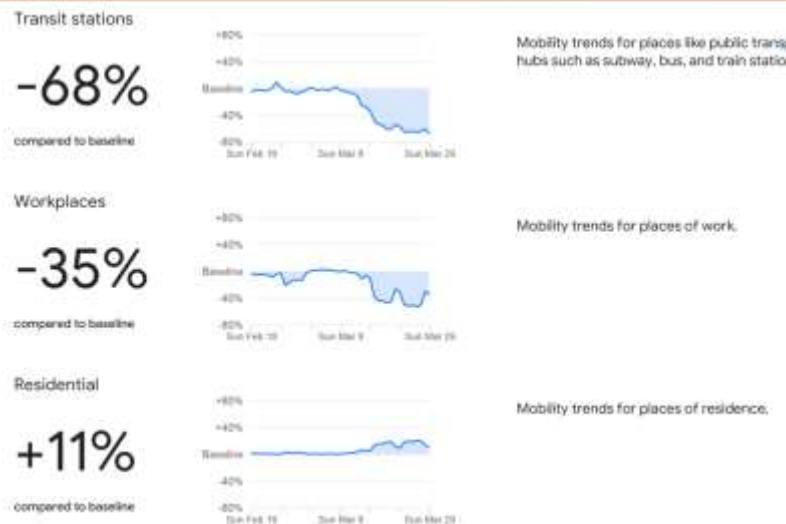


Machine learning-based prediction of COVID-19 diagnosis based on symptoms

Yazeed Zoabi¹, Shira Deri-Rozov¹ and Noam Shomron¹✉

Supervised and unsupervised learning

ADS2



Gedi Lukšys

April 15, 2024



Learning objectives

Today's plan:

- main elements & techniques of supervised learning
- a simple example of supervised learning and gradient ascent approach
- an overview of different supervised and unsupervised learning techniques
- relevant considerations and applications of supervised and unsupervised learning

Supervised learning: what is it for?

Learn to predict relations between 2 sets of data: „inputs“ → „outputs“

Apt. size: 50,
75, 150, 67,
34, 80m²

No of rooms:
2, 4, 5, 3, 2, 3




Rent:
\$700, 950, 1400, 1000,
550, 1200

0, 1, 2, ..., 9

Spam / no Spam

<http://w3.impa.br/~lhf/sib2003/p023/img2.gif>



„Abu Malik:
CONFIDENTIAL TRUST
REPLY URGENTLY“

„bahn.de: Vielen Dank
für Ihren Fahrkartenkauf!
(Auftrag ABCXYZ)“

Input typically multidimensional
(sometimes very much so...)

Output often unidimensional, **discrete**
(classes) or **continuous**

Formally speaking: classification vs. regression

- Supervised learning deals with learning a function $f(X, W) = Y$, based on a number of available data points $(x^{(1)}, y^{(1)})$, $(x^{(2)}, y^{(2)})$, ... , $(x^{(N)}, y^{(N)})$, also called **examples**
 - where $x^{(i)}$ is usually a vector $(x_1^{(i)}, x_2^{(i)}, \dots, x_M^{(i)})$
 - $y^{(i)}$ is either a number (\Rightarrow **regression**)
a class $\{0, 1\}$ or a vector of classes (\Rightarrow **classification**)
- W - **model parameters** (also a vector), M - **dimensionality**
- Most common example: linear regression / classifier

Why is it called „supervised“?

- Suppose our model simulates a function $y = f(x)$
- We can apply it to inputs (x-s) from existing examples to calculate outputs, i.e. $f(x^{(1)})$, $f(x^{(2)})$, ...
- In the beginning the model has not learned to predict the function correctly yet, but as we have the corresponding example („teacher“) outputs, we can measure the mean squared error (MSE):

$$\text{MSE} = ((y^{(1)} - f(x^{(1)}))^2 + \dots + (y^{(n)} - f(x^{(n)}))^2)/n$$

- Sometimes root mean square error (RMSE) is used instead.

This is called the **error or objective** function. The goal is to minimize it.

A simple example - linear model

- Model: $y = ax + b$

Objective function:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - ax^{(i)} - b)^2$$


How can we make our model minimize the error?

We have 2 **parameters**: a & b

- We have to find their values that minimize the error!

How?

1. **Guess!** (need to start somewhere)
2. **Adjust guesses**, e.g. try neighboring values, go in the direction that decreases the error most => **hill climbing**

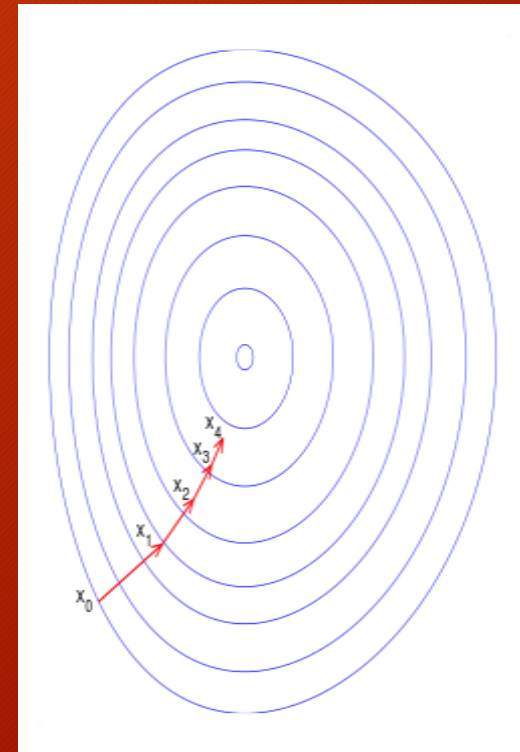


<http://www.statslab.cam.ac.uk/~rrw1/stats/regress1.gif>

Parameter initialization
Typically random

A standard method: gradient descent

- Mean squared error (MSE) is a function of parameters (a and b)
- Derivatives $\delta \text{MSE} / \delta a$ and $\delta \text{MSE} / \delta b$ should be 0 at the minimum of MSE.
- How to get there fastest? (analytical solution aside 😊)
- Go along the steepest direction:
$$\Delta a = -\alpha \delta \text{MSE} / \delta a$$
$$\Delta b = -\alpha \delta \text{MSE} / \delta b$$
- Update of parameters is called a learning rule, and α is the learning rate.



http://en.wikipedia.org/wiki/File:Gradient_descent.svg

Let's try!

$$MSE = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - ax^{(i)} - b)^2$$

- Learning rules:

$$\Delta a = -\alpha / N \sum_{i=1}^N 2(y^{(i)} - ax^{(i)} - b)(-x^{(i)})$$

$$\Delta b = -\alpha / N \sum_{i=1}^N 2(y^{(i)} - ax^{(i)} - b)(-1)$$

- Examples: (0,1), (1,3), (4,6), (2,4), (3,4)
- Try different learning rates

$\alpha = 0.03$

#1: a=0.56, b=0.22
#2: a=0.9, b=0.35
#3: a=1.1, b=0.44
#4: a=1.21, b=0.5
#5: a=1.28, b=0.54
#10: a=1.36, b=0.64
#20: a=1.32, b=0.77
#50: a=1.23, b=1.03
#100: a=1.15, b=1.25
#200: a=1.11, b=1.38

$\alpha = 0.08$

#1: a=1.5, b=0.58
#2: a=1.38, b=0.58
#3: a=1.37, b=0.62
#4: a=1.36, b=0.66
#5: a=1.35, b=0.69
#10: a=1.29, b=0.85
#20: a=1.22, b=1.06
#50: a=1.13, b=1.32
#100: a=1.1, b=1.39
#200: a=1.1, b=1.4

$\alpha = 0.13$

#1: a=2.44, b=0.94
#2: a=0.59, b=0.36
#3: a=1.93, b=0.89
#4: a=0.9, b=0.6
#5: a=1.63, b=0.91
#10: a=1.17, b=0.97
#20: a=1.16, b=1.22
#50: a=1.11, b=1.38
#100: a=1.1, b=1.4
#200: a=1.1, b=1.4

$\alpha = 0.16$


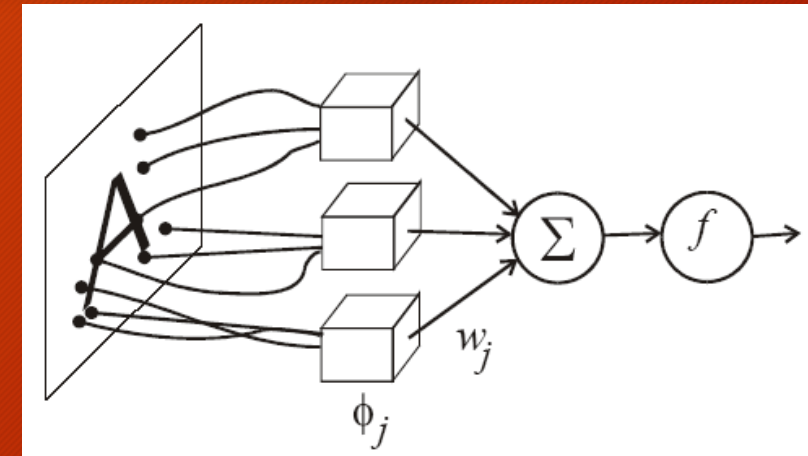
#1: a=3.01, b=1.15
#2: a=-0.5, b=0.01
#3: a=3.46, b=1.48
#4: a=-1.12, b=-0.06
#5: a=4.07, b=1.83
#10: a=-4.25, b=-0.85
#20: a=-19.92, b=-6.11
#50: a=-1206.95, b=-422.37
#100: a=-1030083.04, b=-361332.61
#200: a=-748941284697.65, b=-262714417242.33



Too slow


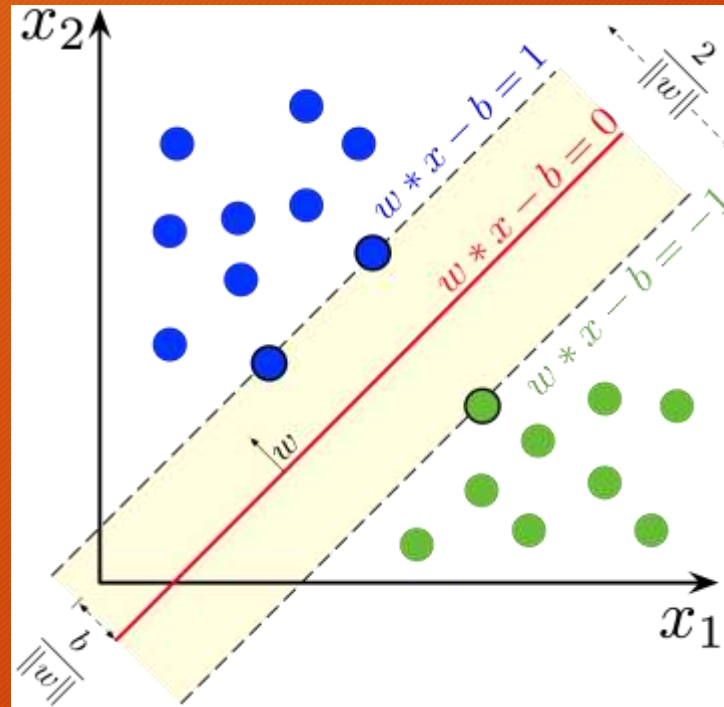
Artificial neural networks (multi-layer perceptrons)

- Inspired by early ideas on neural networks in the brain, particularly the hierarchically organized visual system („the perceptron“)
- Trained using *backpropagation*
- Were shown to have *universal approximation property*
- Still very popular, although recently eclipsed by deep learning (which is a special case)



Support vector machines (SVMs) and kernels

- Finding boundaries with largest margins between classes
- Kernels: projecting to higher dimensions




<http://www.imtech.res.in/raghava/rbpred/svm.jpg>

https://en.wikipedia.org/wiki/Support-vector_machine

Shooting fireworks into the sky or
gene -> protein parallels ☺

Unsupervised learning: a clustering example

- Can be solved with K-means




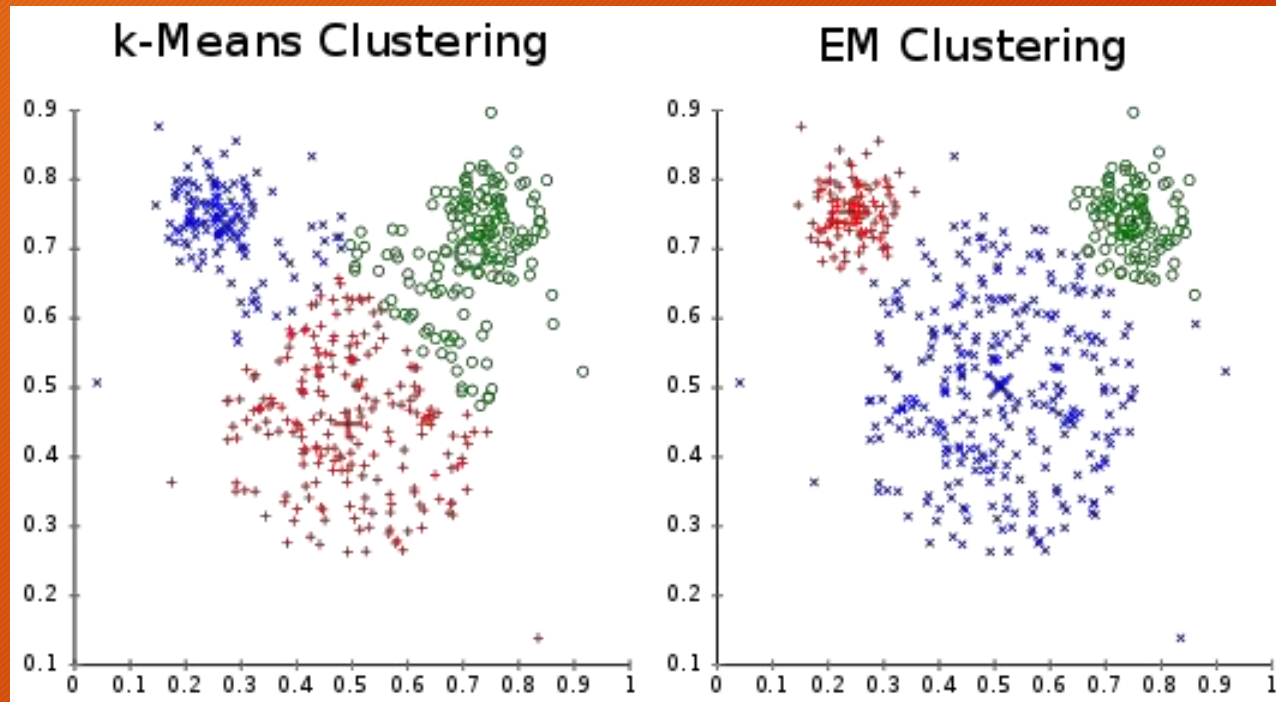
But...

How many clusters?

How to treat borderline points?

More on unsupervised learning: Gaussian mixtures


- Clusters vary in sizes & shapes, modeled by Gaussians
- Calculated using “expectation-maximization” algorithm



Gaussian mixtures for supervised learning

- Classes have complex distributions and may consist of multiple Gaussian clusters

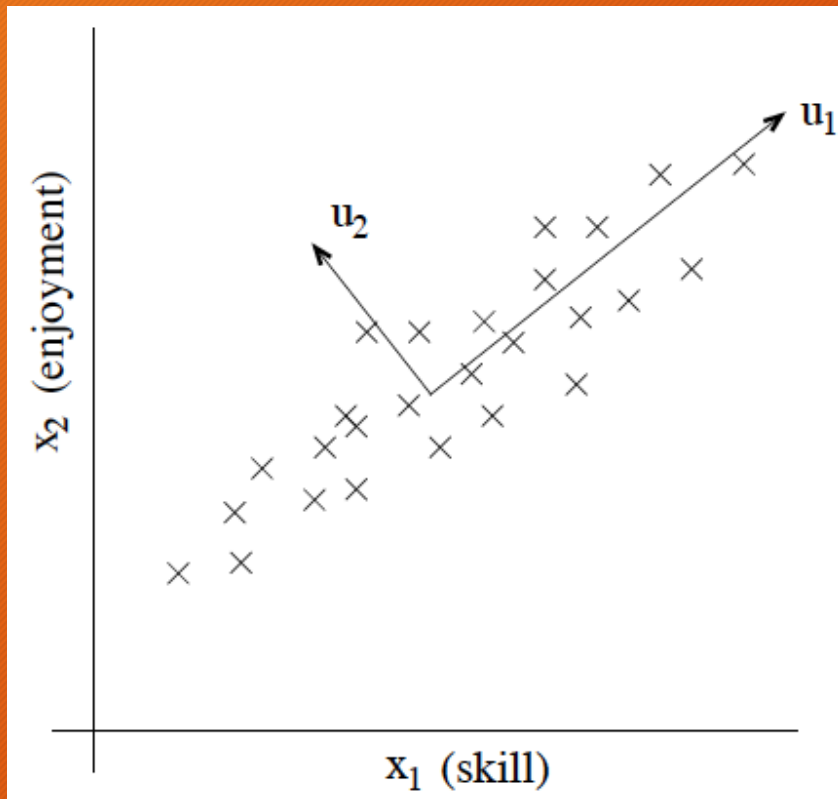
A.W.Moore's bioassay example



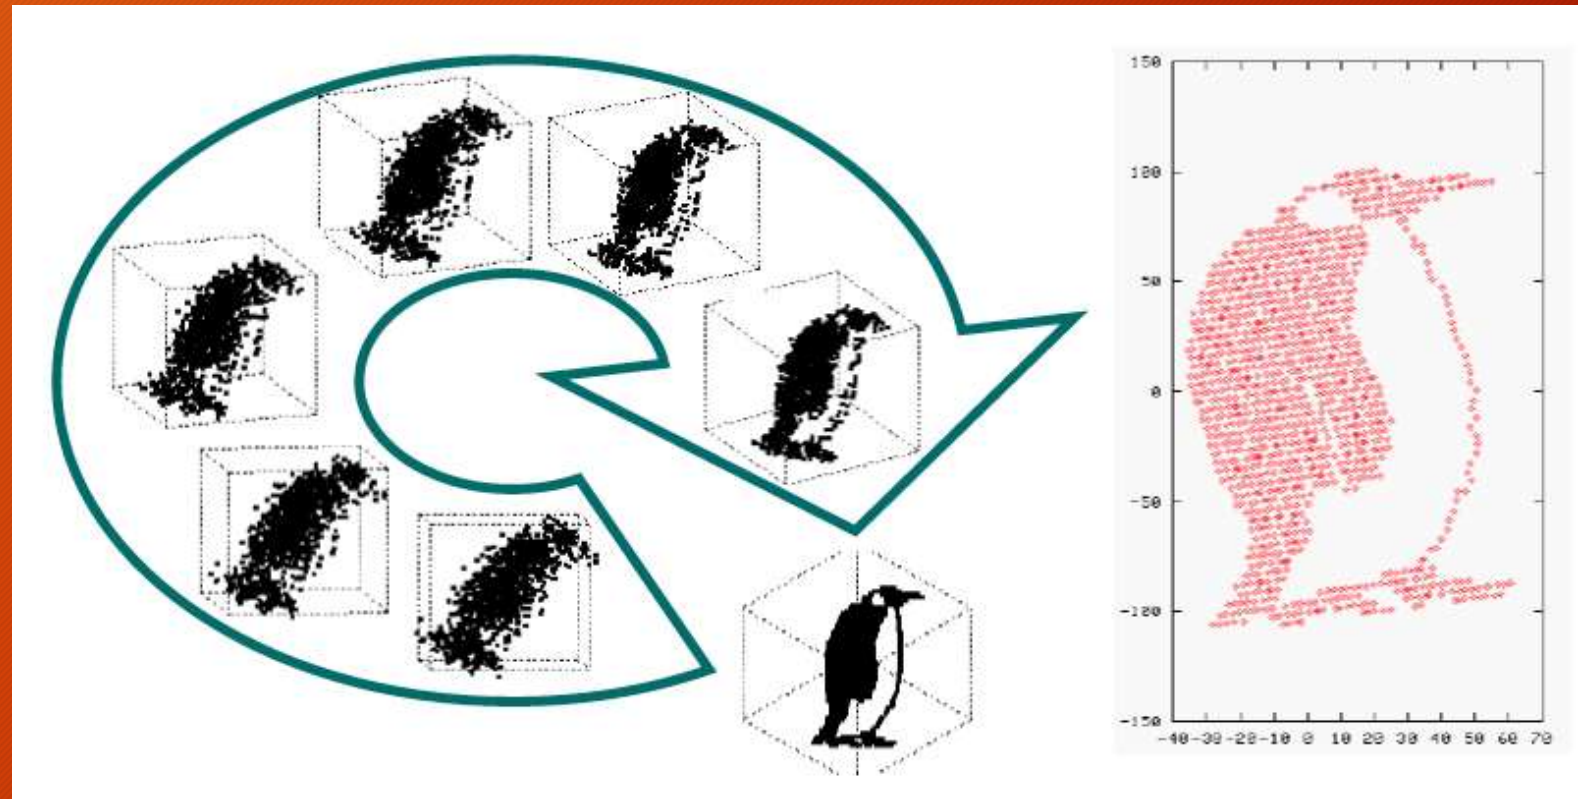
Clustering vs. dimensionality reduction

- Clustering reduces the number of data points (into much fewer clusters)
 - *e.g. K-means, Gaussian mixtures*
- Dimensionality reduction reduces the number of dimensions per data point
 - *e.g. Principal/Independent Component Analyses*
- Use similar principles but in different direction / for different purpose (leading to either fewer data or fewer dimensions)
- In supervised learning dimensionality reduction often used before prediction - **WHY?**

Principal component analysis - rotate data to maximize variance



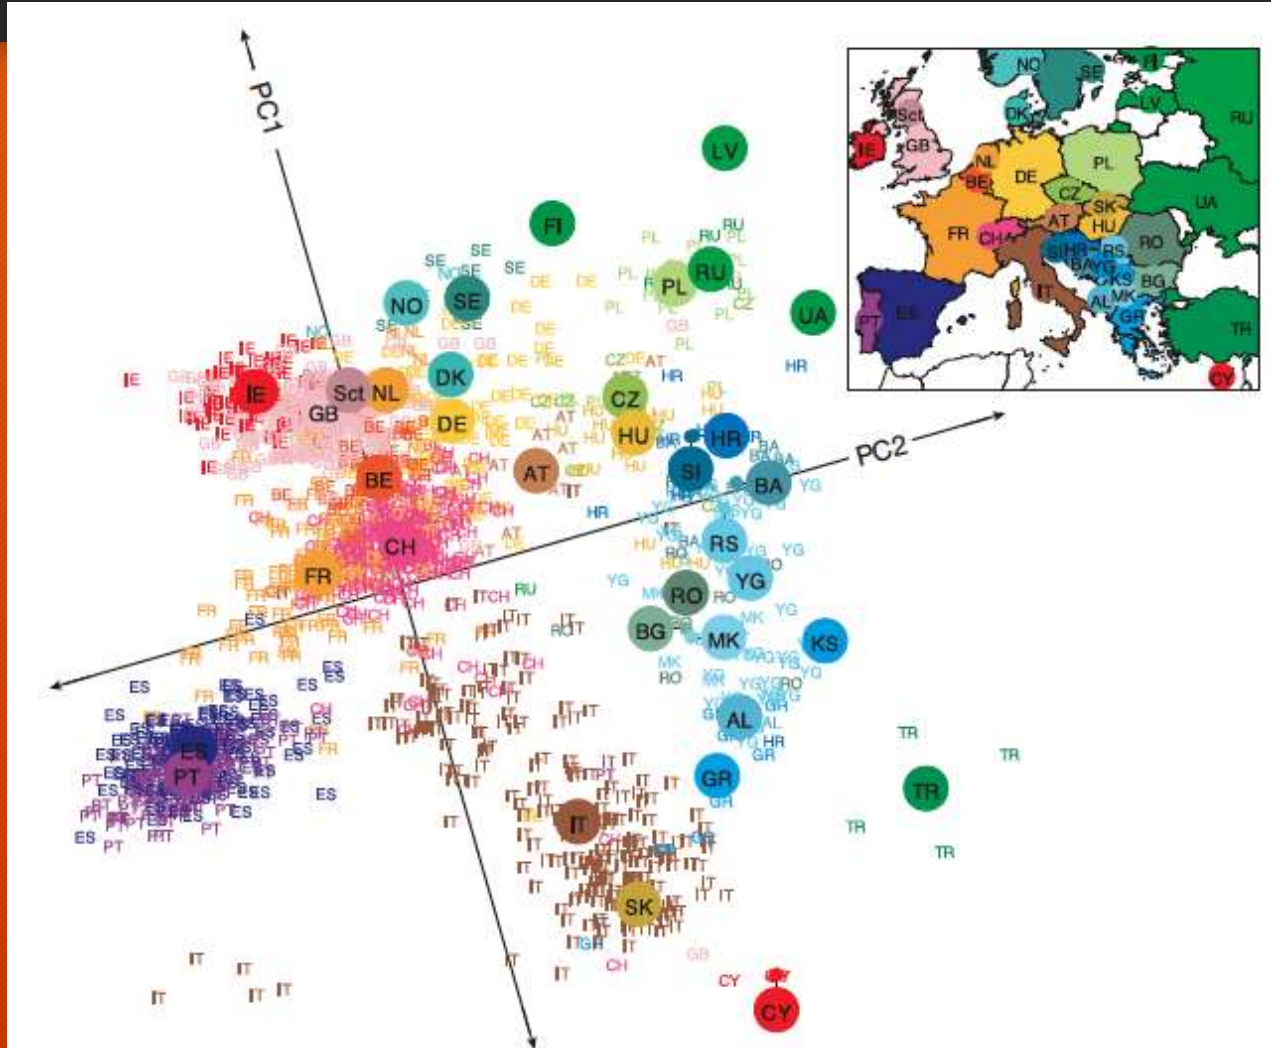
Helicopter driving - Andrew Ng



Example by R.Gutierrez-Osuna

- ## Where they are:

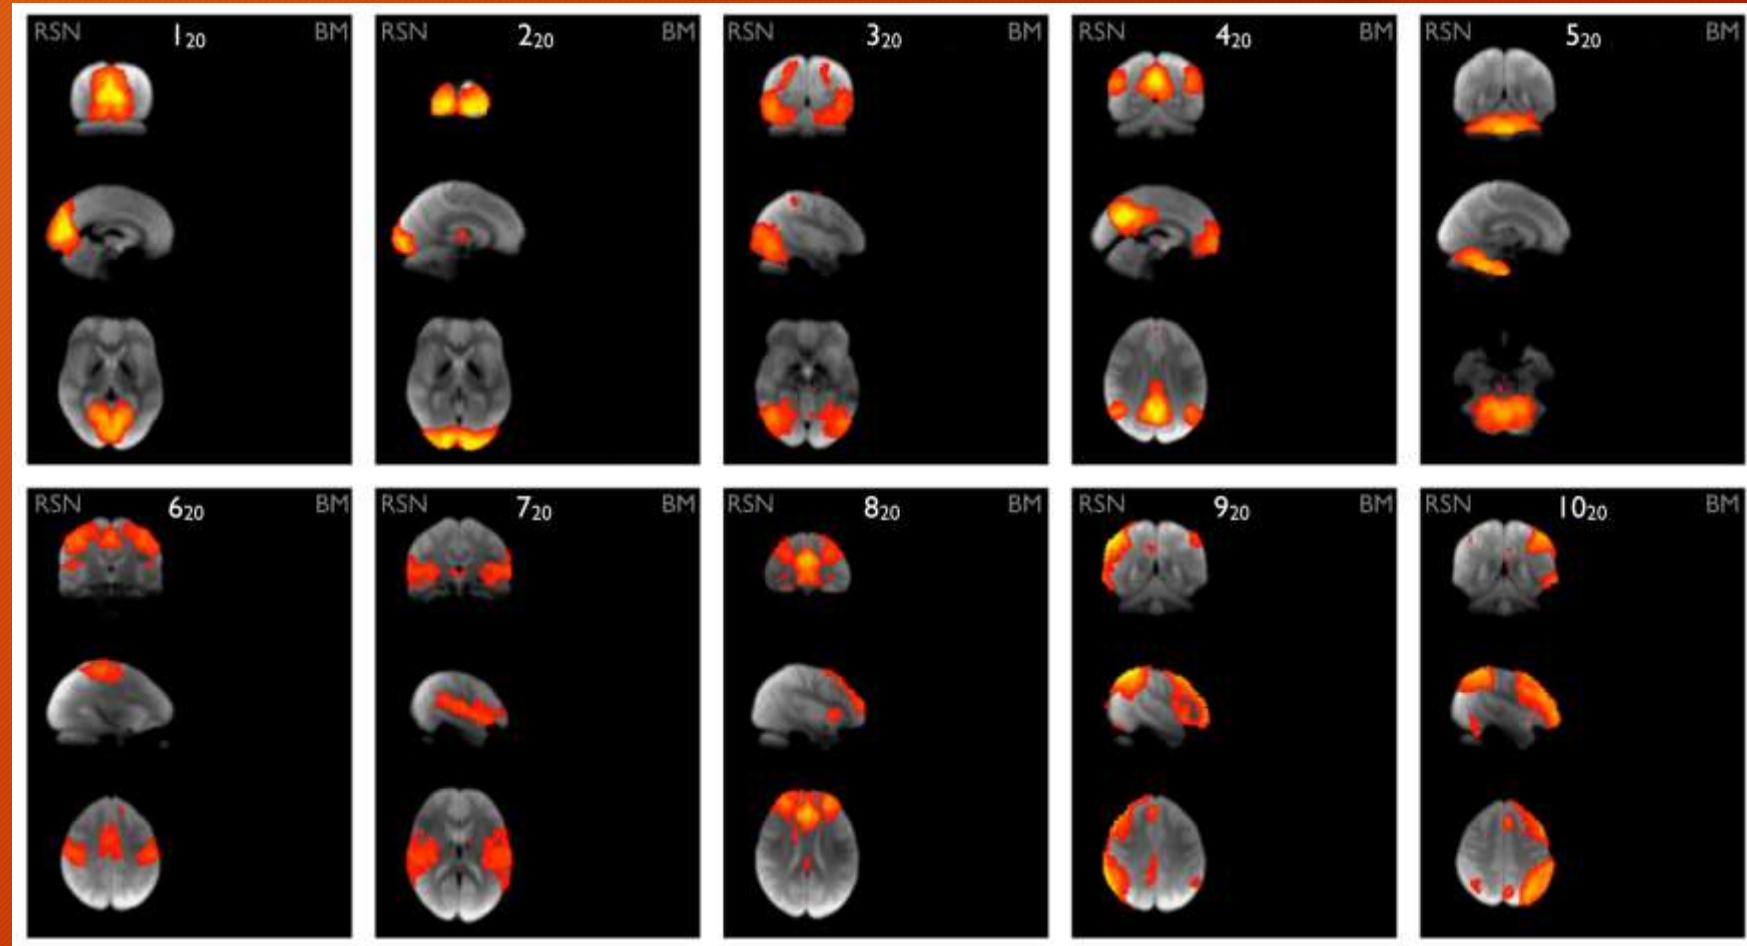

- Novembre
et al.,
Nature 2008



Identification of brain's networks using fMRI + Independent Component Analysis during the resting state

Smith et al.,
PNAS 2009


Independent Component analysis -
extract components
that are not just
uncorrelated but also
independent!



Identification of brain's networks using fMRI + ICA during the resting state & task performance

Extracted components in the resting state and during a variety of task performances are similar!

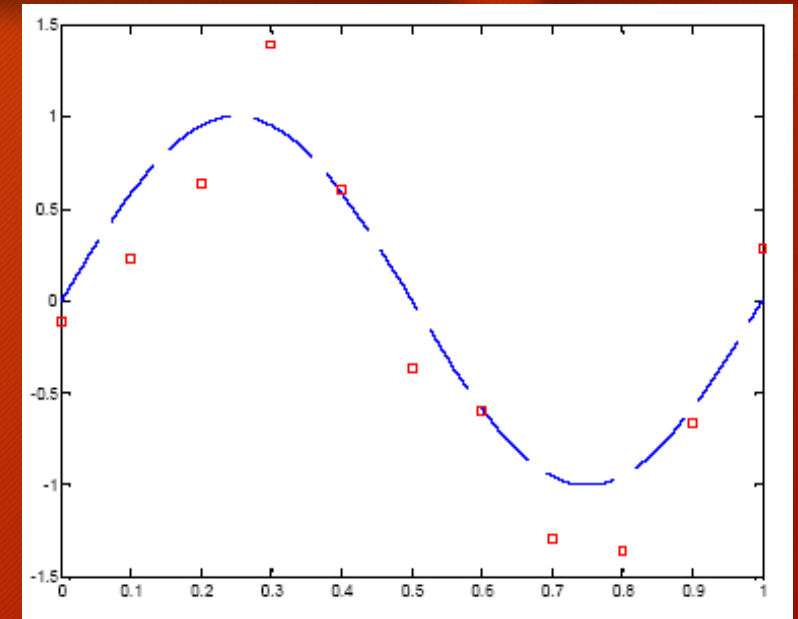
Smith et al., PNAS 2009



Potential problem in supervised learning: overfitting

- Imagine a model: $p(x) = w_0 + w_1x + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j.$
- No. of parameters = $M+1$
- Goal: minimize error

$$MSE_{train} = \frac{1}{N} \sum_{i=1}^N \left(p(x^{(i)}) - y^{(i)} \right)^2$$




Herbert Jaeger, 'Machine Learning', Bremen, 2003

$M=1$, $MSE=0.4852$

$M=3$, $MSE=0.0703$

$M=10$, $MSE = 0!$

Really?



Bias-variance dilemma

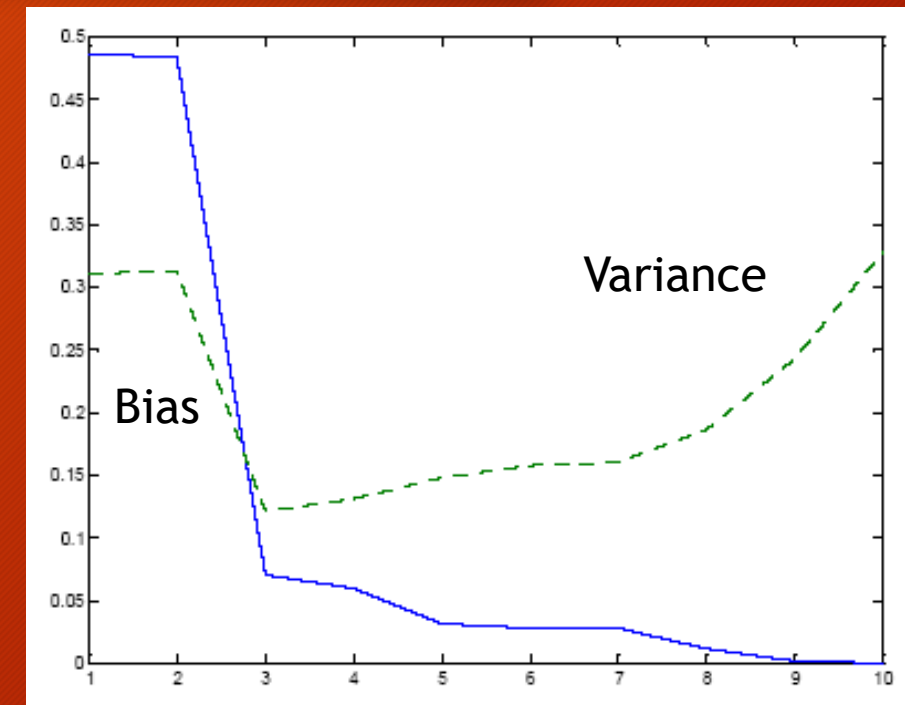
- Why does $M=3$ look better than $M=10$?
 - It better **generalizes** the function! But...
 - We typically don't know what the function is...
- Use the *testing error*!

$$MSE_{test} = \frac{1}{K} \sum_{j=1}^K \left(p(x^{(j)}) - y^{(j)} \right)^2$$

So $M=3$ is really the best 😊

Why „bias-variance“?

- **Bias** - fails to account for relevant data structure
- **Variance** - varies too much depending on training data



H.Jaeger, 'Machine Learning', Bremen, 2003

Dealing with bias-variance dilemma

Finding a right balance (i.e. avoiding overfitting and underfitting) depends on two critical quantities:

- **Model size** (the number of parameters M)
- **The number of samples N**


If $M \geq N$, there will certainly be a substantial overfitting (system of linear equations analogy)

- Normally it's recommended that M would be **much smaller** than N , e.g. 50 or 100 parameters for 1000 samples

How to evaluate over-/underfitting?


- Use cross-validation

- Divide your data into **training** and **validation** sets
- Use the training set for training the model, and the validation set for choosing the number of parameters
- If have too little data, can use „leave-one-out“ cross validation




Dimensionality reduction: features

- So our model size is limited by the number of available samples. **What's the problem with this?**
- Model size is also constrained (*unfortunately from the opposite side*) by the dimensionality of data - usually for each input dimension there is at least 1 parameter



How to resolve this?

- Use not the raw data, but **features!**
- Sometimes this can be done using task-dependent heuristics



<http://w3.impa.br/~lhf/sib2003/p023/img2.gif>